

```
knitr::opts_chunk$set(echo = TRUE)
weather_features <- read.csv("weather_features.csv")
```

## Introduction

weather\_features.csv contains hourly weather data from 2015 to 2018 for 5 major cities in Spain.

```
df <- data.frame(
  Column = colnames(weather_features),
  Description = c("datetime index localized to CET",
    "name of city (Barcelona, Bilbao, Madrid, Seville, Valencia)",
    "temperature (K)",
    "minimum temperature (K)",
    "maximum temperature (K)",
    "pressure (hPa)",
    "humidity (%)",
    "wind speed (m/s)",
    "wind direction (degrees)",
    "rain in last hour (mm)",
    "rain in last 3 hours (mm)",
    "snow in last 3 hours (mm)",
    "cloud cover (%)",
    "weather description - code",
    "weather description - short",
    "weather description - long",
    "weather icon")
)
knitr::kable(df, format = "markdown")
```

Column	Description
dt_iso	datetime index localized to CET
city_name	name of city (Barcelona, Bilbao, Madrid, Seville, Valencia)
temp	temperature (K)
temp_min	minimum temperature (K)
temp_max	maximum temperature (K)
pressure	pressure (hPa)
humidity	humidity (%)
wind_speed	wind speed (m/s)
wind_deg	wind direction (degrees)
rain_1h	rain in last hour (mm)
rain_3h	rain in last 3 hours (mm)
snow_3h	snow in last 3 hours (mm)
clouds_all	cloud cover (%)
weather_id	weather description - code
weather_main	weather description - short
weather_description	weather description - long
weather_icon	weather icon

The last 4 columns are non-numerical, and will be dropped. The data for the city of Barcelona has city\_name " Barcelona", the extra space will be removed.

```
weather_features <- weather_features[1:13]
weather_features[weather_features$city_name == "Barcelona",]$city_name = "Barcelona"
```

## Missing Values

```
sapply(weather_features, function(col)
{ sum(sapply(col, function(x) {
  (is.na(x)) + (x == "")
})) })
```

```
##      dt_iso  city_name      temp  temp_min  temp_max  pressure  humidity
##          0          0          0          0          0          0          0
## wind_speed  wind_deg  rain_1h   rain_3h   snow_3h clouds_all
##          0          0          0          0          0          0
```

Data has no blank or NA cells in any of the columns. Next, we check that there are no missing rows. Since we have hourly data for 3 regular and 1 leap year, we expect there to be  $24 \times 365 \times 4 + 24 = 35064$  unique datetime indices.

```
dates_occurrences <- table(weather_features$dt_iso)
length(dates_occurrences)
```

```
## [1] 35064
```

This matches our expectations.

Since we have 5 cities, we expect each datetime index to appear 5 times.

```
sum(sapply(dates_occurrences, function(date) { date < 5 })))
```

```
## [1] 0
```

This indicates that we have no missing rows.

## Partition Based on City

This means a total of  $35064 \times 5 = 175320$  rows. However, our data has 178396 rows. This indicates the presence of duplicated rows.

To fix this, we first create separate data frames for each city, then remove duplicates.

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 4.4.2
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
## filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
## intersect, setdiff, setequal, union
```

```
Barcelona <- weather_features[weather_features$city_name == "Barcelona", ] %>% distinct()  
Bilbao <- weather_features[weather_features$city_name == "Bilbao", ] %>% distinct()  
Madrid <- weather_features[weather_features$city_name == "Madrid", ] %>% distinct()  
Seville <- weather_features[weather_features$city_name == "Seville", ] %>% distinct()  
Valencia <- weather_features[weather_features$city_name == "Valencia", ] %>% distinct()
```

```
dim(Barcelona)
```

```
## [1] 35064 13
```

```
dim(Bilbao)
```

```
## [1] 35064 13
```

```
dim(Madrid)
```

```
## [1] 35064 13
```

```
dim(Seville)
```

```
## [1] 35064 13
```

```
dim(Valencia)
```

```
## [1] 35064 13
```